

Q2 readme

a. DataLoader Design

- **Standardization:** Images were resized to 224x224 pixels to conform to the input size required by the pre-trained ResNet18 model. Following resizing, we applied normalization using means [0.485, 0.456, 0.406] and standard deviations [0.229, 0.224, 0.225], which are the standard normalization values for models pre-trained on the ImageNet dataset.
- **Label Definition:** Labels were defined based on the filename prefixes, converted to one-hot encoded format to suit the multi-class classification objective.
- **Shuffling:** Data shuffling was employed (shuffle=True) to ensure that each batch seen by the model during training contains a random assortment of all classes, which helps prevent the model from learning spurious patterns.
- **Batch Size:** We chose a batch size of 32, as it is a standard size that balances computational efficiency and model update granularity.

b. Code screenshots related to (a)

```
class WeatherDataset(Dataset):
    def __init__(self, data_dir):
        self.data_dir = data_dir
        self.image_files = os.listdir(data_dir)
        self.transform = transforms.Compose([
            transforms.Resize((224, 224)), # Resize the image to 224x224
            transforms.ToTensor(), # Convert to tensor
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]) # Normalize
        ])

    def __len__(self):
        return len(self.image_files)

    def __getitem__(self, idx):
        image_name = self.image_files[idx]
        image_path = os.path.join(self.data_dir, image_name)
        image = Image.open(image_path).convert('RGB')
        image_tensor = self.transform(image)
        # Extract label information
        label_name = image_name.split('.')[0] # Remove file extension
        label = self.get_label(label_name)

        return image_tensor, label

    def get_label(self, label_name):
        # Return one-hot encoded label based on label name
        if label_name.startswith('Cloudy'):
            return torch.tensor([1, 0, 0, 0, 0])
        elif label_name.startswith('Foggy'):
            return torch.tensor([0, 1, 0, 0, 0])
        elif label_name.startswith('Rainy'):
            return torch.tensor([0, 0, 1, 0, 0])
        elif label_name.startswith('Snowy'):
            return torch.tensor([0, 0, 0, 1, 0])
        elif label_name.startswith('Sunny'):
            return torch.tensor([0, 0, 0, 0, 1])
        else:
            raise ValueError('Invalid label name')
```

```
# Set batch size and shuffle
batch_size = 32
shuffle = True
```

c. A brief introduction to model along with relevant code screenshots

- The core of our classification model is the ResNet18 architecture, renowned for its performance on image data. The model was initialized with weights pre-trained on the ImageNet dataset to leverage the representational learning from a wide and general dataset. We modified the final fully connected layer of the ResNet18 to output five classes, corresponding to our weather classification task: Cloudy, Foggy, Rainy, Snowy, and Sunny.

```
class WeatherClassifier(nn.Module):
    def __init__(self, num_classes):
        super(WeatherClassifier, self).__init__()
        self.base_model = models.resnet18(pretrained=True)
        num_features = self.base_model.fc.in_features
        self.base_model.fc = nn.Linear(num_features, num_classes)

    def forward(self, x):
        x = self.base_model(x)
        return x
```

```
base_model.conv1: Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
base_model.layer1.0.conv1: Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer1.0.conv2: Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer1.1.conv1: Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer1.1.conv2: Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer2.0.conv1: Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
base_model.layer2.0.conv2: Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer2.0.downsample.0: Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
base_model.layer2.1.conv1: Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer2.1.conv2: Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer3.0.conv1: Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
base_model.layer3.0.conv2: Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer3.0.downsample.0: Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
base_model.layer3.1.conv1: Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer3.1.conv2: Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer4.0.conv1: Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
base_model.layer4.0.conv2: Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer4.0.downsample.0: Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
base_model.layer4.1.conv1: Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.layer4.1.conv2: Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
base_model.fc: Linear(in_features=512, out_features=5, bias=True)
Output Shape: [5]
```

- To optimize the model, we used the Adam optimizer with a learning rate of 0.001 and selected the loss function CrossEntropyLoss, which is appropriate for multi-class classification problem.

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

- The model was trained for 50 epochs, and at each epoch, the training loss and accuracy were recorded to monitor the learning progress. And the training and evaluation stages were called separately to ensure the correctness of the model's behavior.

```
# Train the model
for epoch in range(num_epochs):
    model.train() # Make sure the model is in training mode
    running_loss = 0.0
    correct = 0
    total = 0

    for images, labels in dataloader:
        images = images.to(device)
        labels = labels.to(device)

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, torch.argmax(labels, dim=1))

        # Backward pass and optimization
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Calculate accuracy
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == torch.argmax(labels, dim=1)).sum().item()

        running_loss += loss.item()

    # Print loss and accuracy for each epoch
    epoch_loss = running_loss / len(dataloader)
    epoch_acc = correct / total * 100
    print(f"Epoch [{epoch+1}/{num_epochs}]\t Loss: {epoch_loss:.4f}\t Accuracy: {epoch_acc:.2f}%")
```

```
# Evaluate the model on the entire training dataset
model.eval() # Make sure the model is in evaluation mode
total = 0
correct = 0
```

d. Screenshots depicting the training process

Epoch [1/50]	Loss: 0.8685	Accuracy: 67.20%	Epoch [26/50]	Loss: 0.0218	Accuracy: 99.20%
Epoch [2/50]	Loss: 0.5025	Accuracy: 84.40%	Epoch [27/50]	Loss: 0.0315	Accuracy: 98.40%
Epoch [3/50]	Loss: 0.2452	Accuracy: 89.20%	Epoch [28/50]	Loss: 0.0215	Accuracy: 99.20%
Epoch [4/50]	Loss: 0.1419	Accuracy: 94.40%	Epoch [29/50]	Loss: 0.0124	Accuracy: 99.60%
Epoch [5/50]	Loss: 0.1959	Accuracy: 95.20%	Epoch [30/50]	Loss: 0.0127	Accuracy: 99.60%
Epoch [6/50]	Loss: 0.2013	Accuracy: 92.40%	Epoch [31/50]	Loss: 0.0086	Accuracy: 99.60%
Epoch [7/50]	Loss: 0.1565	Accuracy: 94.80%	Epoch [32/50]	Loss: 0.0082	Accuracy: 99.60%
Epoch [8/50]	Loss: 0.0788	Accuracy: 96.80%	Epoch [33/50]	Loss: 0.0094	Accuracy: 99.60%
Epoch [9/50]	Loss: 0.1003	Accuracy: 96.80%	Epoch [34/50]	Loss: 0.0074	Accuracy: 99.60%
Epoch [10/50]	Loss: 0.1160	Accuracy: 96.80%	Epoch [35/50]	Loss: 0.0076	Accuracy: 99.60%
Epoch [11/50]	Loss: 0.1152	Accuracy: 96.00%	Epoch [36/50]	Loss: 0.0068	Accuracy: 99.60%
Epoch [12/50]	Loss: 0.1153	Accuracy: 95.60%	Epoch [37/50]	Loss: 0.0068	Accuracy: 99.20%
Epoch [13/50]	Loss: 0.0863	Accuracy: 97.20%	Epoch [38/50]	Loss: 0.0063	Accuracy: 100.00%
Epoch [14/50]	Loss: 0.0522	Accuracy: 98.00%	Epoch [39/50]	Loss: 0.0062	Accuracy: 99.60%
Epoch [15/50]	Loss: 0.0443	Accuracy: 97.60%	Epoch [40/50]	Loss: 0.0056	Accuracy: 100.00%
Epoch [16/50]	Loss: 0.0192	Accuracy: 99.20%	Epoch [41/50]	Loss: 0.0083	Accuracy: 99.20%
Epoch [17/50]	Loss: 0.1000	Accuracy: 96.80%	Epoch [42/50]	Loss: 0.0058	Accuracy: 99.60%
Epoch [18/50]	Loss: 0.0552	Accuracy: 97.60%	Epoch [43/50]	Loss: 0.0072	Accuracy: 99.60%
Epoch [19/50]	Loss: 0.0469	Accuracy: 98.40%	Epoch [44/50]	Loss: 0.0051	Accuracy: 99.60%
Epoch [20/50]	Loss: 0.0730	Accuracy: 97.60%	Epoch [45/50]	Loss: 0.0073	Accuracy: 99.60%
Epoch [21/50]	Loss: 0.0684	Accuracy: 98.00%	Epoch [46/50]	Loss: 0.0087	Accuracy: 99.60%
Epoch [22/50]	Loss: 0.0536	Accuracy: 97.60%	Epoch [47/50]	Loss: 0.0070	Accuracy: 99.20%
Epoch [23/50]	Loss: 0.0682	Accuracy: 97.20%	Epoch [48/50]	Loss: 0.0093	Accuracy: 99.60%
Epoch [24/50]	Loss: 0.0663	Accuracy: 97.20%	Epoch [49/50]	Loss: 0.0047	Accuracy: 100.00%
Epoch [25/50]	Loss: 0.0616	Accuracy: 98.00%	Epoch [50/50]	Loss: 0.0052	Accuracy: 99.60%

e. Accuracy on the training set (with screenshots)

Accuracy on training set: 99.60%

References:

[1] ResNet-18 实现 Cifar-10 图像分类 <https://blog.csdn.net/sunqiande88/article/details/80100891>